

# Laboratory Work 8: SQL Indexes

## Objective

To understand and practice creating and managing different types of indexes in PostgreSQL, learning how indexes improve query performance and when to use them effectively.

## Prerequisites

- Completion of Laboratory Work 7 (Views and Roles)
- Basic SQL knowledge (SELECT, INSERT, UPDATE, DELETE)
- Understanding of query performance concepts
- Access to a PostgreSQL database system

## Theoretical Background

### What are Indexes?

Indexes are database objects that improve the speed of data retrieval operations on tables. They work like an index in a book - instead of reading every page to find information, you can look up the index to quickly locate what you need.

### Why Use Indexes?

- **Speed up queries:** Faster data retrieval, especially on large tables
- **Improve JOIN performance:** Speed up queries that join multiple tables
- **Enforce uniqueness:** Ensure unique values in columns
- **Sort data efficiently:** Help with ORDER BY operations

### Index Types in PostgreSQL

1. **B-tree:** Default type, good for equality and range queries
2. **Hash:** Only for equality comparisons
3. **GiST:** For geometric and full-text search data
4. **GIN:** For array and full-text search operations

### Important Considerations

- Indexes speed up reads but slow down writes (INSERT, UPDATE, DELETE)
- Don't create indexes on every column - only on frequently queried columns
- Indexes consume disk space
- PostgreSQL automatically creates indexes for PRIMARY KEY and UNIQUE constraints

## Part 1: Database Setup

Use the same tables from **Lab 6** and **Lab 7** (employees, departments, projects). If you need to recreate them, here's a quick setup:

```

sql
-- Create tables
CREATE TABLE departments (
    dept_id INT PRIMARY KEY,
    dept_name VARCHAR(50),
    location VARCHAR(50)
);

CREATE TABLE employees (
    emp_id INT PRIMARY KEY,
    emp_name VARCHAR(100),
    dept_id INT,
    salary DECIMAL(10,2),
    FOREIGN KEY (dept_id) REFERENCES departments(dept_id)
);

CREATE TABLE projects (
    proj_id INT PRIMARY KEY,
    proj_name VARCHAR(100),
    budget DECIMAL(12,2),
    dept_id INT,
    FOREIGN KEY (dept_id) REFERENCES departments(dept_id)
);

```

-- Insert sample data

```

INSERT INTO departments VALUES
(101, 'IT', 'Building A'),
(102, 'HR', 'Building B'),
(103, 'Operations', 'Building C');

INSERT INTO employees VALUES
(1, 'John Smith', 101, 50000),
(2, 'Jane Doe', 101, 55000),
(3, 'Mike Johnson', 102, 48000),
(4, 'Sarah Williams', 102, 52000),
(5, 'Tom Brown', 103, 60000);

INSERT INTO projects VALUES
(201, 'Website Redesign', 75000, 101),
(202, 'Database Migration', 120000, 101),
(203, 'HR System Upgrade', 50000, 102);

```

## Part 2: Creating Basic Indexes

### Exercise 2.1: Create a Simple B-tree Index

Create an index on the `salary` column of the `employees` table:

sql

```
CREATE INDEX emp_salary_idx ON employees(salary);
```

**Verify the index was created:**

sql

-- List all indexes on employees table

```
SELECT indexname, indexdef
```

```
FROM pg_indexes
```

```
WHERE tablename = 'employees';
```

**Question:** How many indexes exist on the employees table? (Hint: PRIMARY KEY creates an automatic index)

## Exercise 2.2: Create an Index on a Foreign Key

Create an index on the dept\_id column in the employees table:

sql

```
CREATE INDEX emp_dept_idx ON employees(dept_id);
```

**Test the index usage:**

sql

-- This query should use the index

```
SELECT * FROM employees WHERE dept_id = 101;
```

**Question:** Why is it beneficial to index foreign key columns?

## Exercise 2.3: View Index Information

Use PostgreSQL system catalogs to view detailed index information:

sql

-- View all indexes in your database

```
SELECT
```

```
tablename,
```

```
indexname,
```

```
indexdef
```

```
FROM pg_indexes
```

```
WHERE schemaname = 'public'
```

```
ORDER BY tablename, indexname;
```

**Question:** List all the indexes you see. Which ones were created automatically?

## Part 3: Multicolumn Indexes

## Exercise 3.1: Create a Multicolumn Index

Create an index on both `dept_id` and `salary` columns:

sql

```
CREATE INDEX emp_dept_salary_idx ON employees(dept_id, salary);
```

**Test the multicolumn index:**

sql

-- This query can use the multicolumn index

```
SELECT emp_name, salary  
FROM employees  
WHERE dept_id = 101 AND salary > 52000;
```

**Question:** Would this index be useful for a query that only filters by `salary` (without `dept_id`)? Why or why not?

## Exercise 3.2: Understanding Column Order

Create another multicolumn index with reversed column order:

sql

```
CREATE INDEX emp_salary_dept_idx ON employees(salary, dept_id);
```

**Compare with queries:**

sql

-- Query 1: Filters by `dept_id` first

```
SELECT * FROM employees WHERE dept_id = 102 AND salary > 50000;
```

-- Query 2: Filters by `salary` first

```
SELECT * FROM employees WHERE salary > 50000 AND dept_id = 102;
```

**Question:** Does the order of columns in a multicolumn index matter? Explain.

## Part 4: Unique Indexes

### Exercise 4.1: Create a Unique Index

First, add a new column for employee email:

sql

```
ALTER TABLE employees ADD COLUMN email VARCHAR(100);  
  
UPDATE employees SET email = 'john.smith@company.com' WHERE emp_id = 1;  
UPDATE employees SET email = 'jane.doe@company.com' WHERE emp_id = 2;  
UPDATE employees SET email = 'mike.johnson@company.com' WHERE emp_id = 3;  
UPDATE employees SET email = 'sarah.williams@company.com' WHERE emp_id = 4;  
UPDATE employees SET email = 'tom.brown@company.com' WHERE emp_id = 5;
```

Now create a unique index on the email column:

```
sql  
CREATE UNIQUE INDEX emp_email_unique_idx ON employees(email);  
Test the uniqueness constraint:
```

```
sql  
-- This should fail with a unique violation error  
INSERT INTO employees (emp_id, emp_name, dept_id, salary, email)  
VALUES (6, 'New Employee', 101, 55000, 'john.smith@company.com');  
Question: What error message did you receive?
```

## Exercise 4.2: Unique Index vs UNIQUE Constraint

Check what indexes exist after adding a UNIQUE constraint:

```
sql  
-- Add a phone column with UNIQUE constraint  
ALTER TABLE employees ADD COLUMN phone VARCHAR(20) UNIQUE;  
View the indexes:
```

```
sql  
SELECT indexname, indexdef  
FROM pg_indexes  
WHERE tablename = 'employees' AND indexname LIKE '%phone%';  
Question: Did PostgreSQL automatically create an index? What type of index?
```

## Part 5: Indexes and Sorting

### Exercise 5.1: Create an Index for Sorting

Create an index optimized for descending salary queries:

```
sql  
CREATE INDEX emp_salary_desc_idx ON employees(salary DESC);  
Test with an ORDER BY query:
```

```
sql  
SELECT emp_name, salary  
FROM employees  
ORDER BY salary DESC;
```

**Question:** How does this index help with ORDER BY queries?

## Exercise 5.2: Index with NULL Handling

Create an index that handles NULL values specially:

```
sql  
CREATE INDEX proj_budget_nulls_first_idx ON projects(budget NULLS FIRST);  
Test the index:
```

```
sql  
SELECT proj_name, budget  
FROM projects  
ORDER BY budget NULLS FIRST;
```

## Part 6: Indexes on Expressions

### Exercise 6.1: Create a Function-Based Index

Create an index for case-insensitive employee name searches:

```
sql  
CREATE INDEX emp_name_lower_idx ON employees(LOWER(emp_name));  
Test the expression index:
```

```
sql  
-- This query can use the expression index  
SELECT * FROM employees WHERE LOWER(emp_name) = 'john smith';  
Question: Without this index, how would PostgreSQL search for names case-insensitively?
```

### Exercise 6.2: Index on Calculated Values

Add a hire\_date column and create an index on the year:

sql

```
ALTER TABLE employees ADD COLUMN hire_date DATE;
```

```
UPDATE employees SET hire_date = '2020-01-15' WHERE emp_id = 1;
UPDATE employees SET hire_date = '2019-06-20' WHERE emp_id = 2;
UPDATE employees SET hire_date = '2021-03-10' WHERE emp_id = 3;
UPDATE employees SET hire_date = '2020-11-05' WHERE emp_id = 4;
UPDATE employees SET hire_date = '2018-08-25' WHERE emp_id = 5;
```

-- Create index on the year extracted from hire\_date

```
CREATE INDEX emp_hire_year_idx ON employees(EXTRACT(YEAR FROM hire_date));
```

**Test the index:**

sql

```
SELECT emp_name, hire_date
FROM employees
WHERE EXTRACT(YEAR FROM hire_date) = 2020;
```

## Part 7: Managing Indexes

### Exercise 7.1: Rename an Index

Rename the emp\_salary\_idx index to employees\_salary\_index:

sql

```
ALTER INDEX emp_salary_idx RENAME TO employees_salary_index;
```

**Verify the rename:**

sql

```
SELECT indexname FROM pg_indexes WHERE tablename = 'employees';
```

### Exercise 7.2: Drop Unused Indexes

Drop the redundant multicolumn index we created earlier:

sql

```
DROP INDEX emp_salary_dept_idx;
```

**Question:** Why might you want to drop an index?

## Exercise 7.3: Reindex

Rebuild an index to optimize its structure:

```
sql  
REINDEX INDEX employees_salary_index;
```

**When is REINDEX useful?**

- After bulk INSERT operations
- When index becomes bloated
- After significant data modifications

## Part 8: Practical Scenarios

### Exercise 8.1: Optimize a Slow Query

Consider this query that runs frequently:

```
sql  
SELECT e.emp_name, e.salary, d.dept_name  
FROM employees e  
JOIN departments d ON e.dept_id = d.dept_id  
WHERE e.salary > 50000  
ORDER BY e.salary DESC;
```

**Create indexes to optimize this query:**

```
sql  
-- Index for the WHERE clause  
CREATE INDEX emp_salary_filter_idx ON employees(salary) WHERE salary > 50000;  
  
-- Index for the JOIN  
-- (already created: emp_dept_idx)  
  
-- Index for ORDER BY  
-- (already created: emp_salary_desc_idx)
```

### Exercise 8.2: Partial Index

Create an index only for high-budget projects (budget > 80000):

```
sql  
CREATE INDEX proj_high_budget_idx ON projects(budget)  
WHERE budget > 80000;
```

## Test the partial index:

```
sql
SELECT proj_name, budget
FROM projects
WHERE budget > 80000;
```

**Question:** What's the advantage of a partial index compared to a regular index?

## Exercise 8.3: Analyze Index Usage

Use EXPLAIN to see if indexes are being used:

```
sql
EXPLAIN SELECT * FROM employees WHERE salary > 52000;
```

**Question:** Does the output show an "Index Scan" or a "Seq Scan" (Sequential Scan)? What does this tell you?

## Part 9: Index Types Comparison

### Exercise 9.1: Create a Hash Index

Create a hash index on department name:

```
sql
CREATE INDEX dept_name_hash_idx ON departments USING HASH (dept_name);
```

**Test the hash index:**

```
sql
SELECT * FROM departments WHERE dept_name = 'IT';
```

**Question:** When should you use a HASH index instead of a B-tree index?

### Exercise 9.2: Compare Index Types

Create both B-tree and Hash indexes on the project name:

```
sql
-- B-tree index
CREATE INDEX proj_name_btree_idx ON projects(proj_name);
```

```
-- Hash index  
CREATE INDEX proj_name_hash_idx ON projects USING HASH (proj_name);  
Test with different queries:
```

```
sql  
-- Equality search (both can be used)  
SELECT * FROM projects WHERE proj_name = 'Website Redesign';  
  
-- Range search (only B-tree can be used)  
SELECT * FROM projects WHERE proj_name > 'Database';
```

## Part 10: Cleanup and Best Practices

### Exercise 10.1: Review All Indexes

List all indexes and their sizes:

```
sql  
SELECT  
    schemaname,  
    tablename,  
    indexname,  
    pg_size.pretty(pg_relation_size(indexname::regclass)) AS index_size  
FROM pg_indexes  
WHERE schemaname = 'public'  
ORDER BY tablename, indexname;  
Question: Which index is the largest? Why?
```

### Exercise 10.2: Drop Unnecessary Indexes

Identify and drop indexes that are duplicates or rarely used:

```
sql  
-- Drop the duplicate expression indexes  
DROP INDEX IF EXISTS proj_name_hash_idx;
```

-- Keep only necessary indexes

### Exercise 10.3: Document Your Indexes

Create a view that documents all custom indexes:

```

sql
CREATE VIEW index_documentation AS
SELECT
    tablename,
    indexname,
    indexdef,
    'Improves salary-based queries' as purpose
FROM pg_indexes
WHERE schemaname = 'public'
AND indexname LIKE '%salary%';

SELECT * FROM index_documentation;

```

## Summary Questions

1. What is the default index type in PostgreSQL?
2. Name three scenarios where you should create an index:
3. Name two scenarios where you should NOT create an index:
4. What happens to indexes when you INSERT, UPDATE, or DELETE data?
5. How can you check if a query is using an index?

## Best Practices Checklist

- Index columns used frequently in WHERE clauses
- Index foreign key columns
- Index columns used in JOIN conditions
- Index columns used in ORDER BY
- Don't over-index (indexes have overhead)
- Consider multicolumn indexes for queries with multiple filters
- Use partial indexes for frequently queried subsets
- Regularly analyze and remove unused indexes
- Use EXPLAIN to verify index usage
- Consider expression indexes for computed values

## Additional Challenges (Optional)

1. Create an index that would optimize finding all employees hired in a specific month
2. Create a composite unique index on dept\_id and email in employees table
3. Use EXPLAIN ANALYZE to compare query performance with and without indexes
4. Create a covering index that includes all columns needed for a specific query

### Submission Requirements:

- SQL script with all CREATE INDEX statements
- Answers to all questions
- Screenshots of EXPLAIN output for at least one query
- Brief reflection on when to use indexes

Good luck!